

Transitioning from DP to RL

Chaitra Gopalappa



References

- **Sutton and Barto, Second edition**
 - Chapter 2, Section 2.4 and 2.5
 - Chapter 3
 - Chapter 6, Section 6.4 and 6.5

- A factory is interested in determining whether to order parts (raw material) or not based on inventory at the end of each week. Below are the details of this factory
 - N products manufactured;
 - Demand for each $i \sim \text{Poisson}(\lambda_i \text{ per week})$
 - Maximum backorder = B_i
 - Revenue per product = r_i
 - Manufacturing these products requires a total of M number of parts
 - Maximum inventory capacity for each part $j = k_j$
 - Total inventory capacity for all parts = K
 - Shipping and ordering costs for each part = S_j
 - Delivery times from supplier for each part $\sim \text{Normal}(\mu_j, \sigma_j)$
 - Not all parts are needed for all products
 - Parts assembled in multiple steps
 - Assembly time for step k of product $i \sim \text{Normal}(\bar{\mu}_{ik}, \bar{\sigma}_{ik})$

Challenges with DP: #1

- Curse of modeling
 - Generating TPM and TRewardM
 - Transitions are a function of multiple random variables
 - Difficult to estimate complex pdfs
- Solution: maximum likelihood estimation through historical data or simulation generated data
 - $P(i, a, j) \approx \frac{w(i, a, j)}{V(i, a)}$
 - $V(i, a)$ = number of time action a taken when system is in state i
 - $w(i, a, j)$ = number of times transition was to state j if action a taken when system was in state i
- If MLE feasible use DP: as DP guarantees optimality
 - Example: in factory problem in slide 1, $N=2$, $M=2$ (when problem is small)

Challenges with DP: #2

- In factory problem in slide 1, if N =tens or hundreds, M =hundreds or thousands?
- Curse of dimensionality
 - Large number of states
 - Estimation issue: # of samples needed to estimate TPM and TRM will be large
- Solution:
 - Reinforcement learning: method of performing DP within a simulator

Solution algorithms to solve ‘discrete time’ MDP

- DP policy iteration
 - DP value iteration
 - Model building algorithms: Build TPM and TRM using simulation and then apply DP
 - Q-learning
 - DQL: deep learning + RL
 - Policy gradient
- Dynamic programming
(model-based)
- Reinforcement learning
(model-free space)

Theory behind RL: Transitioning DP to RL (discounted total reward)

1. Bellman optimality equation to derive **Q-factors**
 1. $Q(s,a)$: Action-value functions (recollect we used state-value functions $V(s)$ in DP)
2. Q-factor version of Bellman optimality equation
3. Q-factor value iteration – still DP domain- but updating Q-factors instead of value function
4. Robbins- Monro algorithm – to estimate mean of RV from sample
5. Using Robbins-Monro to update Q-factors- leads to ‘model-free’ algorithm
6. Finally, leads to Q-learning value iteration algorithm

Each step will be discussed next

Bellman equation for a “fixed” (deterministic) policy π

$$V(s) = \sum_{s'} p(s, \pi(s), s') [r + \lambda V(s')]$$

- $r = r(s, \pi(s), s')$; r here is deterministic
- Note: if there is randomness in value of r we rewrite the value function as
 - $V(s) = \sum_{s', r} [p(s, \pi(s), s', r) [r + \lambda V(s')]]$

State-value
function

Bellman “optimality” equation

$$v^*(s) = \max_{a \in A} \sum_{s'} [p(s, a, s') \cdot (r(s, a, s') + \lambda v^*(s'))]$$

- v^* is the optimal value function for the MDP. r here is deterministic
- Note: if there is randomness in value of r we rewrite
 - $v^*(s) = \max_{a \in A} \sum_{s', r} [p(s', r | s, a) \cdot (r(s, a, s') + \lambda v^*(s'))]$
- Furthermore, it is the only function satisfying the property

Value function

$$v^*(s) = \max_{a \in A} \sum_{s'} p(s, a, s') [r(s, a, s') + \lambda v^*(s')]$$

State-value
function

Q-factor version of Bellman optimality (Chapter 3, Sutton and Barto)

- $Q(s, a) = \sum_{s'} p(s, a, s') [r(s, a, s') + \lambda v^*(s')]$
- $v^*(s) = \max_{a \in A} Q(s, a)$
- Combining the two:
- $Q(s, a) = \sum_{s'} p(s, a, s') \left[r(s, a, s') + \lambda \max_{a'} Q(s', a') \right]$

Action-value
function

Q-factor version of value iteration (still in DP)

1. Initialize

1. $Q(s, a) \in \mathbb{R}$, arbitrarily $\forall (s, a) \in (\mathcal{S}, A)$;
2. Set $V(s) = \max_{a \in A} Q(s, a), \forall s \in \mathcal{S}$
3. set tolerance $\theta (=1e-6)$

2. For each (s, a) pair compute Q-factor:

$$Q(s, a) = \sum_{s'} p(s, a, s') \left[r(s, a, s') + \lambda \max_{a'} Q(s', a') \right]$$

3. For each s compute value function

$$v \leftarrow V(s)$$

$$V(s) = \max_{a \in A} Q(s, a)$$

If $(|v - V(s)|) < \theta$ goto step 4, else goto step 2

4. Find corresponding optimal policy

$$- \pi(s) = \operatorname{argmax}_a Q(s, a)$$

Except for this additional step to calculate Q-values, everything else is the same as in value iteration

Robbins-Monro algorithm

- Suppose X is a random variable, with x_i is the i^{th} independent sample of X
- Then, $\mathbb{E}[X] = \lim_{k \rightarrow \infty} \frac{\sum_{i=1}^k x_i}{k}$
- Robbins-Monro algorithm utilizes above to allow for incremental updating (**Sutton and Barto, Section 2.4 and 2.5**):
 - Let $X_k = \frac{\sum_{i=1}^k x_i}{k}$ then $X_{k+1} = \frac{\sum_{i=1}^{k+1} x_i}{k+1} = \frac{\sum_{i=1}^k x_i + x_{k+1}}{k+1}$
 - $X_{k+1} = \frac{(kX_k + X_k - X_k + x_{k+1})}{k+1} = \frac{X_k(k+1) - X_k + x_{k+1}}{k+1}$
 - $X_{k+1} = \frac{X_k(k+1)}{k+1} - \frac{X_k}{k+1} + \frac{x_{k+1}}{k+1} = X_k - \frac{X_k}{k+1} + \frac{x_{k+1}}{k+1}$
 - $X_{k+1} = X_k + \frac{1}{k+1} [x_{k+1} - X_k] \sim X_k + \alpha [x_{k+1} - X_k]$
 - $NewEstimate \leftarrow OldEstimate + StepSize [Sample - OldEstimate]$; $StepSize = \alpha = \frac{1}{k+1}$

Robbins-Monro to Q-factor updating

$$\begin{aligned} \text{R-M: } X_{k+1} &= X_k + \alpha[x_{k+1} - X_k] \\ X_{k+1} &= X_k + \alpha[\text{sample} - X_k] \end{aligned}$$

- $Q(s, a) = \sum_{s'} p(s, a, s') \left[r(s, a, s') + \lambda \max_{a'} Q(s', a') \right]$
- $\sim Q(s, a) = \mathbb{E} \left[r(s, a, s') + \lambda \max_{a'} Q(s', a') \right]$
- *Applying Robbins-Monro*
 - $Q(s, a)$ is random variable, then $\left[r(s, a, s') + \lambda \max_{a'} Q(s', a') \right]$ is a sample of that random variable
 - $Q(s, a) = Q(s, a) + \alpha[\text{sample} - Q(s, a)]$
 - $Q(s, a) = Q(s, a) + \alpha \left[r(s, a, s') + \lambda \max_{a'} Q(s', a') - Q(s, a) \right]$
 - $Q_{k+1}(S_t, A_t) = Q_k(S_t, A_t) + \alpha \left[R_{t+1} + \lambda \max_{a'} Q(S_{t+1}, a') - Q_k(S_t, A_t) \right]$
 - $Q(S_t, A_t) = Q(S_t, A_t) + \alpha \left[R_{t+1} + \lambda \max_{a'} Q(S_{t+1}, a') - Q(S_t, A_t) \right]$

Robbins-Monro to Q-factor updating

$$\begin{aligned} \text{R-M: } X_{k+1} &= X_k + \alpha [x_{k+1} - X_k] \\ X_{k+1} &= X_k + \alpha [\text{sample} - X_k] \end{aligned}$$

- $Q(s, a) = \sum_{s'} p(s, a, s') [r(s, a, s') + \lambda \max_{a'} Q(s', a')]$
- $\sim Q(s, a) = \mathbb{E} [r(s, a, s') + \lambda \max_{a'} Q(s', a')] \sim \mathbb{E}[\text{sample}]$
- *Applying Robbins-Monro*
 - $Q(s, a)$ is random variable, then $[r(s, a, s') + \lambda \max_{a'} Q(s', a')]$ is a sample of that random variable
 - $Q(s, a) = Q(s, a) + \alpha [\text{sample} - Q(s, a)]$
 - $Q(s, a) = Q(s, a) + \alpha [r(s, a, s') + \lambda \max_{a'} Q(s', a') - Q(s, a)]$
 - $Q(S_t, A_t) = Q(S_t, A_t) + \alpha [R_{t+1} + \lambda \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$
- $\text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize} [\text{Target (or sample)} - \text{OldEstimate}] ; \alpha = \frac{1}{k+1}$

Notice there is no p (no TPM) \Rightarrow model free!

- Note: **sample** is more of a **target** here given the dynamics
- $[\text{Target} - \text{OldEstimate}]$ is the error in the estimate, which can be reduced by taking a step toward the target
- If we set $\alpha = \frac{1}{k+1}$ it is direct averaging of samples, however, this can be altered
- Generally, convergence with probability 1 if following conditions are satisfied,
 - $\sum_{k=1}^{\infty} \alpha_k = \infty$ and $\sum_{k=1}^{\infty} \alpha_k^2 < \infty$
- Or, specific to action selection, suppose $\alpha_n(a)$ is step-size after n^{th} selection of action a
 - $\sum_{n=1}^{\infty} \alpha_n(a) = \infty$ and $\sum_{n=1}^{\infty} \alpha_n^2(a) < \infty$

Recollect these are the same as stochastic approximation convergence properties when we did gradient descent; In fact, 'incremental' stochastic gradient descent is sometimes attributed to R-M

UMassAmherst
The Commonwealth's Flagship Campus